

Brain Stroke Prediction

Importing Libraries

```
In [2]: import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import matplotlib.pyplot as plt #creating static, interactive, and animated visualizations in Python
import seaborn as sns #data visualizations
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
from sklearn.model_selection import train_test_split
from sklearn.pipeline import Pipeline, make_pipeline
from sklearn.utils.validation import check_is_fitted

import warnings
warnings.filterwarnings('ignore')
```

Viewing the data in pandas

```
In [3]: #The dataframe is read from the csv file
df=pd.read_csv('brain_stroke.csv')
df.head(10)
```

```
Out[3]:
```

	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_glucose
0	Male	67.0	0	1	Yes	Private	Urban	
1	Male	80.0	0	1	Yes	Private	Rural	
2	Female	49.0	0	0	Yes	Private	Urban	
3	Female	79.0	1	0	Yes	Self-employed	Rural	
4	Male	81.0	0	0	Yes	Private	Urban	
5	Male	74.0	1	1	Yes	Private	Rural	
6	Female	69.0	0	0	No	Private	Urban	
7	Female	78.0	0	0	Yes	Private	Urban	
8	Female	81.0	1	0	Yes	Private	Rural	
9	Female	61.0	0	1	Yes	Govt_job	Rural	

```
In [4]: df.tail(5)
```

Out[4]:

	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_glu
--	--------	-----	--------------	---------------	--------------	-----------	----------------	---------

4976	Male	41.0	0	0	No	Private	Rural	
4977	Male	40.0	0	0	Yes	Private	Urban	
4978	Female	45.0	1	0	Yes	Govt_job	Rural	
4979	Male	40.0	0	0	Yes	Private	Rural	
4980	Female	80.0	1	0	Yes	Private	Urban	

In [5]: *#Count of number of records and columns*
df.shape

Out[5]: (4981, 11)

In [6]: *# Provides the data type of all attributes and the number of NOT NULL values count is obtained*
df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4981 entries, 0 to 4980
Data columns (total 11 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   gender          4981 non-null   object
1   age             4981 non-null   float64
2   hypertension     4981 non-null   int64
3   heart_disease   4981 non-null   int64
4   ever_married    4981 non-null   object
5   work_type       4981 non-null   object
6   Residence_type  4981 non-null   object
7   avg_glucose_level 4981 non-null   float64
8   bmi             4981 non-null   float64
9   smoking_status  4981 non-null   object
10  stroke          4981 non-null   int64
dtypes: float64(3), int64(3), object(5)
memory usage: 428.2+ KB
```

In [7]: df.describe()

Out[7]:

	age	hypertension	heart_disease	avg_glucose_level	bmi	stroke
count	4981.000000	4981.000000	4981.000000	4981.000000	4981.000000	4981.000000
mean	43.419859	0.096165	0.055210	105.943562	28.498173	0.049789
std	22.662755	0.294848	0.228412	45.075373	6.790464	0.217531
min	0.080000	0.000000	0.000000	55.120000	14.000000	0.000000
25%	25.000000	0.000000	0.000000	77.230000	23.700000	0.000000
50%	45.000000	0.000000	0.000000	91.850000	28.100000	0.000000
75%	61.000000	0.000000	0.000000	113.860000	32.600000	0.000000
max	82.000000	1.000000	1.000000	271.740000	48.900000	1.000000

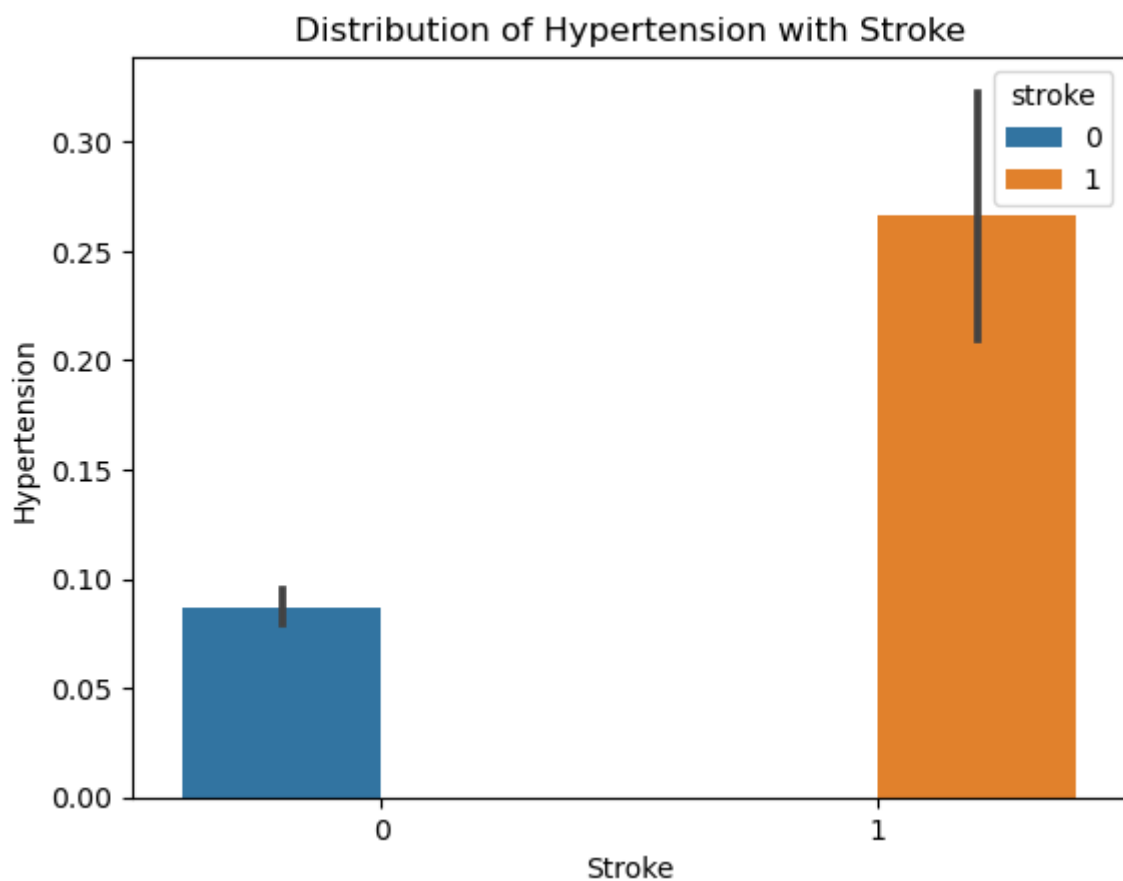
In [8]: df.isnull().sum()

```
Out[8]: gender      o
age      o
hypertension  o
heart_disease  o
ever_married  o
work_type  o
Residence_type  o
avg_glucose_level  o
bmi      o
smoking_status  o
stroke    o
dtype: int64
```


```
In [9]: # Checking the values in the stroke column
print(df["stroke"].value_counts())
```

```
0    4733
1     248
Name: stroke, dtype: int64
```

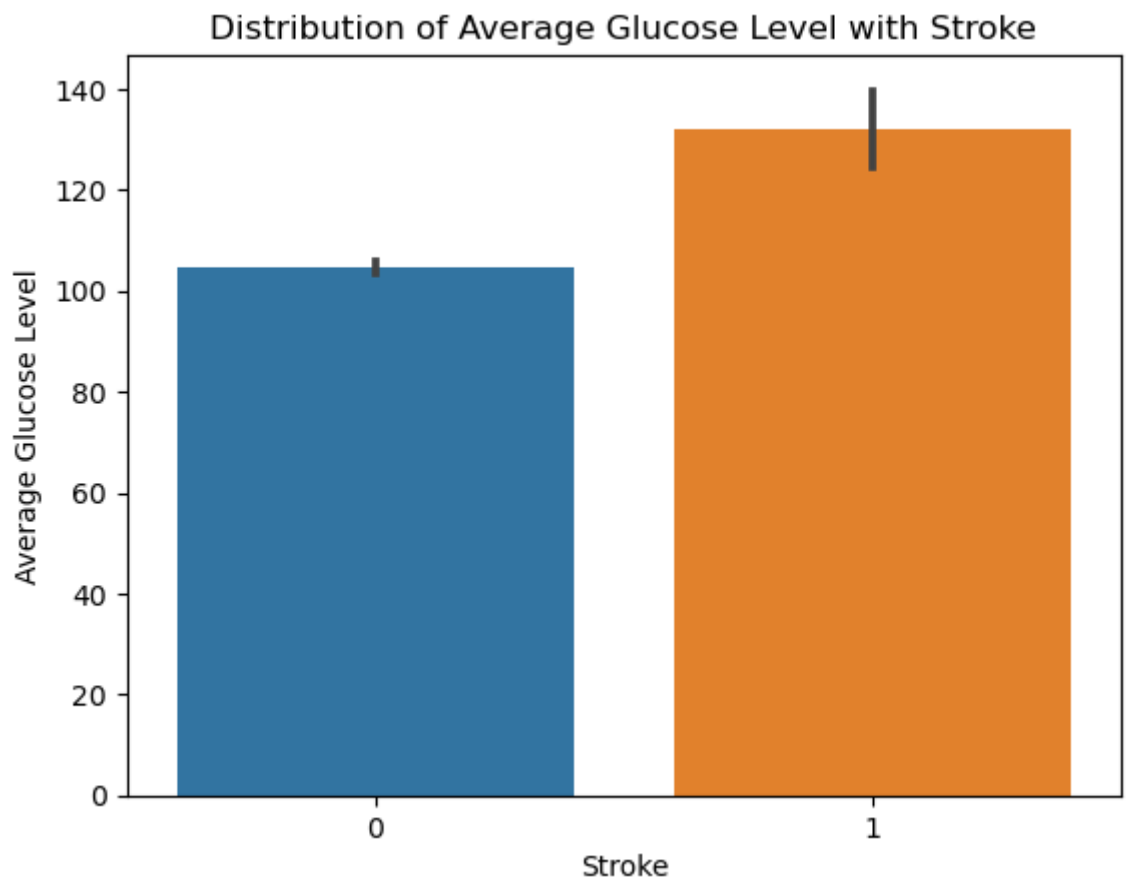
```
In [11]: sns.barplot(x="stroke", y='hypertension', data=df, hue="stroke")
#label axis
plt.xlabel("Stroke")
plt.ylabel("Hypertension")
plt.title("Distribution of Hypertension with Stroke");
```



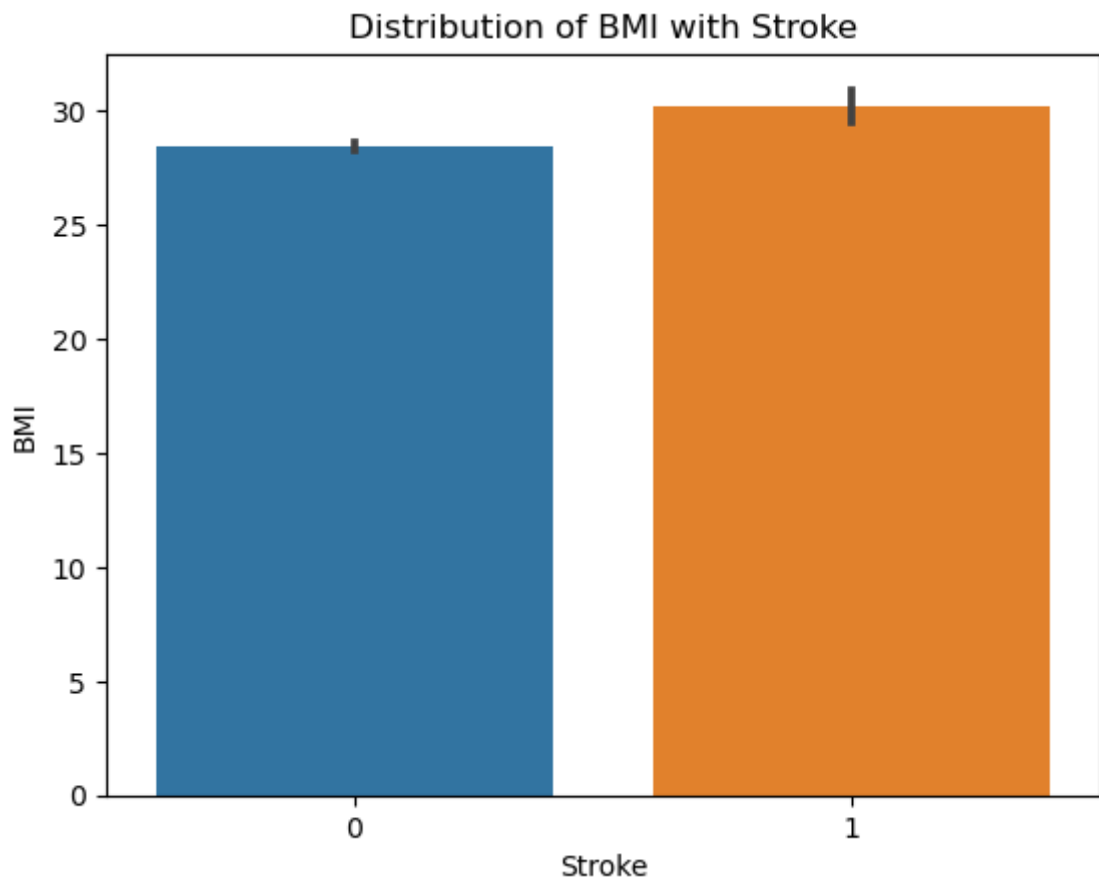
```
In [12]: sns.barplot(x="stroke", y='heart_disease', data=df)
#label axis
plt.xlabel("Stroke")
plt.ylabel("Heart Disease")
plt.title("Distribution of Heart Disease with Stroke");
```



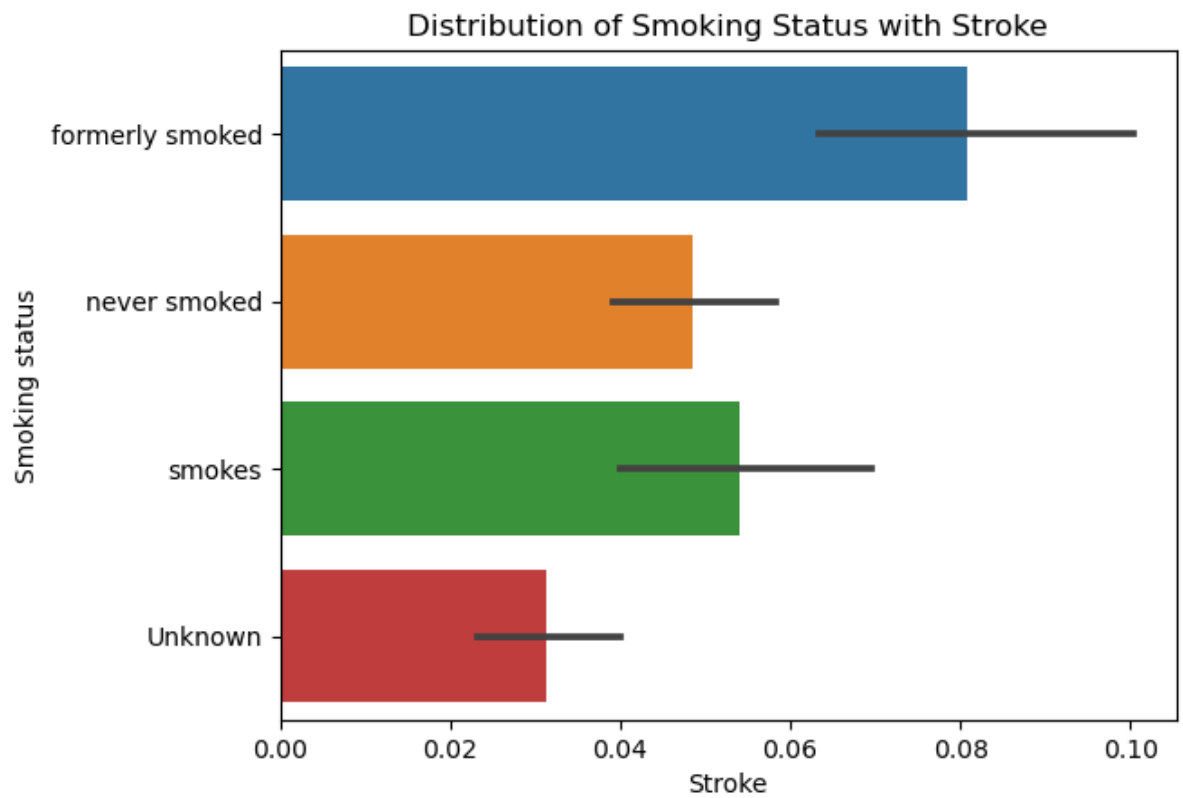
```
In [13]: sns.barplot(x="stroke", y='avg_glucose_level', data=df)
#label axis
plt.xlabel("Stroke")
plt.ylabel("Average Glucose Level")
plt.title("Distribution of Average Glucose Level with Stroke");
```




```
In [14]: sns.barplot(x="stroke", y='bmi', data=df)
#label axis
plt.xlabel("Stroke")
plt.ylabel("BMI")
plt.title("Distribution of BMI with Stroke");
```



```
In [15]: sns.barplot(x="stroke", y='smoking_status', data=df)
#label axis
plt.xlabel("Stroke")
plt.ylabel("Smoking status")
plt.title("Distribution of Smoking Status with Stroke");
```



```
In [16]: sns.barplot(x="stroke", y='age', data=df)
#label axis
plt.xlabel("Stroke")
plt.ylabel("Age")
plt.title("Distribution of Age with Stroke");
```



```
In [17]: print(df["smoking_status"].value_counts())
```

```
never smoked    1838
Unknown         1500
formerly smoked  867
smokes          776
Name: smoking_status, dtype: int64
```

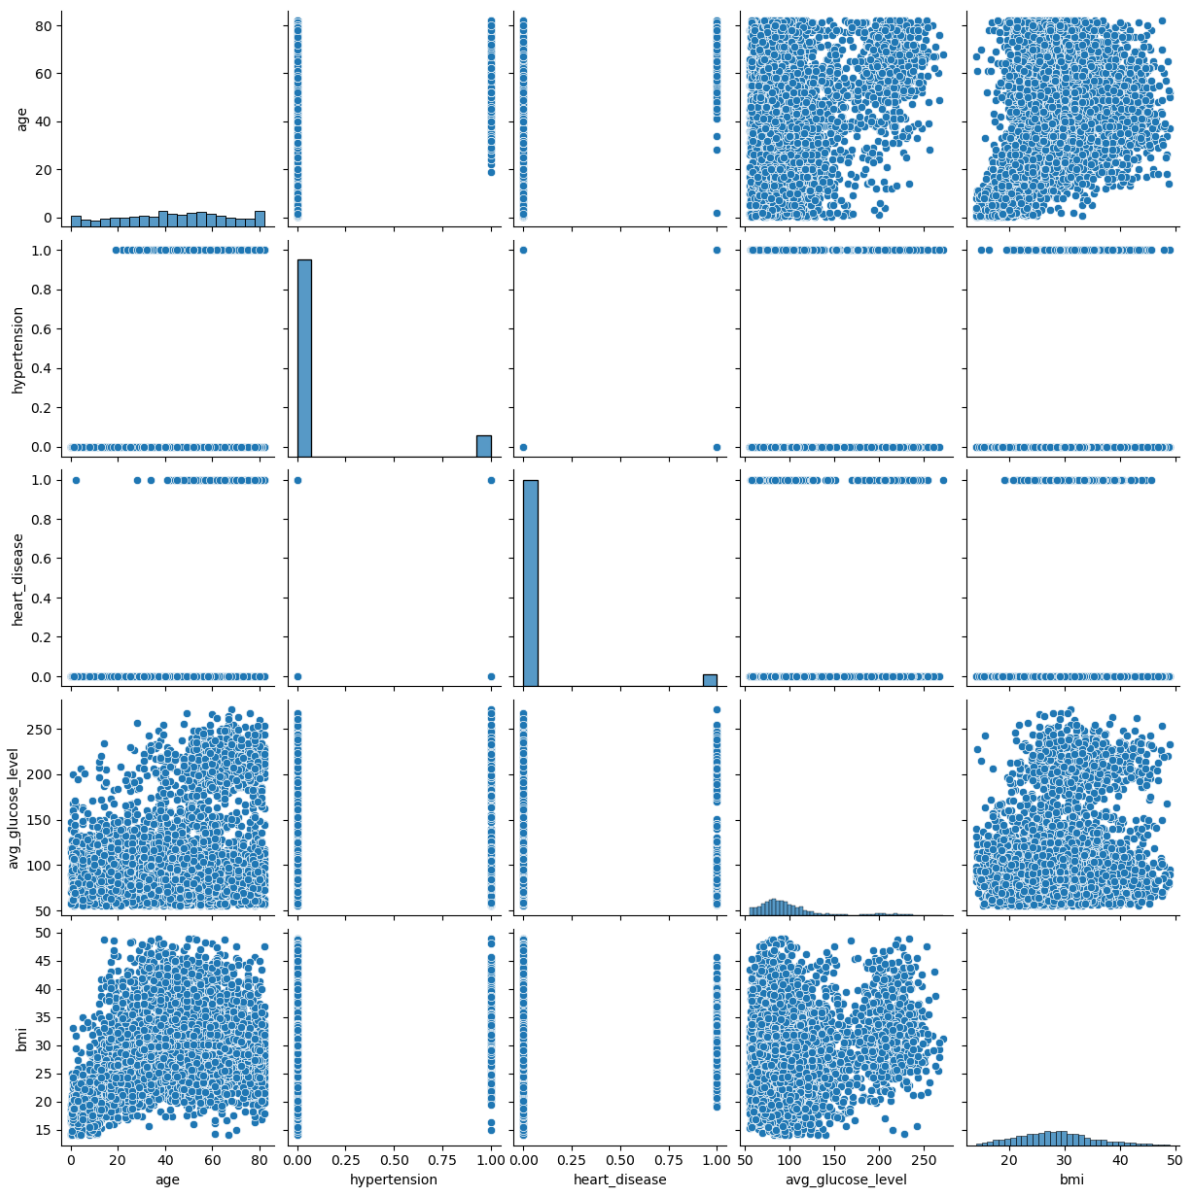
```
In [18]: print(df["heart_disease"].value_counts())
```

```
0    4706
1      275
Name: heart_disease, dtype: int64
```

```
In [19]: # Excluding the object datatype in the dataset
df = df.select_dtypes(exclude=['object'])
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4981 entries, 0 to 4980
Data columns (total 6 columns):
#   Column          Non-Null Count  Dtype
---  -
0   age             4981 non-null   float64
1   hypertension     4981 non-null   int64
2   heart_disease    4981 non-null   int64
3   avg_glucose_level 4981 non-null   float64
4   bmi              4981 non-null   float64
5   stroke           4981 non-null   int64
dtypes: float64(3), int64(3)
memory usage: 233.6 KB
```

```
In [20]: y = df["stroke"].values
x_data = df.drop(["stroke"], axis = 1)
sns.pairplot(x_data)
plt.show()
```



Data Normalization

```
In [21]: x = (x_data - np.min(x_data))/(np.max(x_data) - np.min(x_data)).values
x.head()
```

```
Out[21]:
```

	age	hypertension	heart_disease	avg_glucose_level	bmi
0	0.816895	0.0	1.0	0.801265	0.647564
1	0.975586	0.0	1.0	0.234512	0.530086
2	0.597168	0.0	0.0	0.536008	0.584527
3	0.963379	1.0	0.0	0.549349	0.286533
4	0.987793	0.0	0.0	0.605161	0.429799

Correlation Matrix

```
In [22]: df.corr() #calculates the correlation matrix between the numerical columns
```


Out[22]:


	age	hypertension	heart_disease	avg_glucose_level	bmi	stroke
age	1.000000	0.278120	0.264852	0.236763	0.373703	0.246478
hypertension	0.278120	1.000000	0.111974	0.170028	0.158762	0.131965
heart_disease	0.264852	0.111974	1.000000	0.166847	0.060926	0.134610
avg_glucose_level	0.236763	0.170028	0.166847	1.000000	0.186348	0.133227
bmi	0.373703	0.158762	0.060926	0.186348	1.000000	0.056926
stroke	0.246478	0.131965	0.134610	0.133227	0.056926	1.000000

In [23]:

```
# Correlation matrix between the attributes in the dataset to find if any attributes are correlated
corrmat=df.corr()
f,ax=plt.subplots(figsize=(9,8))
sns.heatmap(corrmat,ax=ax,cmap="YlGnBu",linewidth=0.8,annot=True)
```

Out[23]:

<AxesSubplot:>



Splitting the dataset

In [24]:

```
target = "stroke"
X = df.drop(columns=target)
y = df[target]
```

20% test set & 80% train set

```
In [25]: X_train, X_test, y_train, y_test = train_test_split(
        X, y, test_size=0.2, random_state=42
    )
```

```
print("X_train shape:", X_train.shape)
print("y_train shape:", y_train.shape)
print("X_test shape:", X_test.shape)
print("y_test shape:", y_test.shape)
```

```
X_train shape: (3984, 5)
y_train shape: (3984,)
X_test shape: (997, 5)
y_test shape: (997,)
```

```
In [26]: acc_baseline = y_train.value_counts(normalize=True).max()
        print("Baseline Accuracy:", round(acc_baseline, 2))
```

```
Baseline Accuracy: 0.95
```

Building Model

```
In [27]: # Build model using Logistic Regression
        model = LogisticRegression(max_iter=1000)
```

```
# Fit model to training data
model.fit(X_train, y_train)
```

```
Out[27]: LogisticRegression(max_iter=1000)
```

```
In [28]: # Predicting the result
        y_pred = model.predict(X_test)
```

```
In [30]: model.score(X_test, y_test)
```

```
Out[30]: 0.9458375125376128
```


```
In [31]: # Confusion Matrix
        from sklearn.metrics import accuracy_score, confusion_matrix, classification_report, plot_confusion_matrix
        ac = accuracy_score(y_test, y_pred) * 100
        cr = classification_report(y_test, y_pred)
        print("accuracy score : ", ac)
        print("classification report :", cr)
```

```
accuracy score : 94.58375125376128
classification report :
```

	0	0.95	1.00	0.97	943
	1	0.00	0.00	0.00	54
accuracy				0.95	997
macro avg		0.47	0.50	0.49	997
weighted avg		0.89	0.95	0.92	997

```
In [32]: plot_confusion_matrix(model, X_test, y_test)
```

```
Out[32]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x1e212924b20>
```



In []: